

1.

$$\theta^T X^T = [\theta_0 + \theta_1 x_1^{(1)} + \dots + \theta_n x_n^{(1)}, \dots, \theta_0 + \theta_1 x_1^{(m)} + \dots + \theta_n x_n^{(m)}]$$

$$X \cdot \theta = \begin{bmatrix} 1 & x_1^{(1)} & \dots & x_n^{(1)} \\ \vdots & \vdots & & \vdots \\ 1 & x_1^{(m)} & \dots & x_n^{(m)} \end{bmatrix} \cdot \begin{bmatrix} \theta_0 \\ \vdots \\ \theta_n \end{bmatrix} = \begin{bmatrix} \theta_0 + \theta_1 x_1^{(1)} + \dots + \theta_n x_n^{(1)} \\ \vdots \\ \theta_0 + \theta_1 x_1^{(m)} + \dots + \theta_n x_n^{(m)} \end{bmatrix}$$

$$\Rightarrow \theta^T X^T X \theta = [\theta_0 + \theta_1 x_1^{(1)} + \dots + \theta_n x_n^{(1)}]^2 + \dots + [\theta_0 + \theta_1 x_1^{(m)} + \dots + \theta_n x_n^{(m)}]^2$$

$$\frac{\partial \theta^T X^T X \theta}{\partial \theta} = 2 \cdot \begin{bmatrix} (\theta_0 + \theta_1 x_1^{(1)} + \dots + \theta_n x_n^{(1)}) + \dots + (\theta_0 + \theta_1 x_1^{(m)} + \dots + \theta_n x_n^{(m)}) \\ x_1^{(1)} \cdot (\theta_0 + \theta_1 x_1^{(1)} + \dots + \theta_n x_n^{(1)}) + \dots + x_1^{(m)} \cdot (\theta_0 + \theta_1 x_1^{(m)} + \dots + \theta_n x_n^{(m)}) \\ \vdots \\ x_n^{(1)} \cdot (\theta_0 + \theta_1 x_1^{(1)} + \dots + \theta_n x_n^{(1)}) + \dots + x_n^{(m)} \cdot (\theta_0 + \theta_1 x_1^{(m)} + \dots + \theta_n x_n^{(m)}) \end{bmatrix}$$

Meanwhile,

$$X^T \cdot X \theta = \begin{bmatrix} 1 & \dots & 1 \\ x_1^{(1)} & \dots & x_1^{(m)} \\ \vdots & & \vdots \\ x_n^{(1)} & \dots & x_n^{(m)} \end{bmatrix} \cdot \begin{bmatrix} \theta_0 + \theta_1 x_1^{(1)} + \dots + \theta_n x_n^{(1)} \\ \vdots \\ \theta_0 + \theta_1 x_1^{(m)} + \dots + \theta_n x_n^{(m)} \end{bmatrix}$$

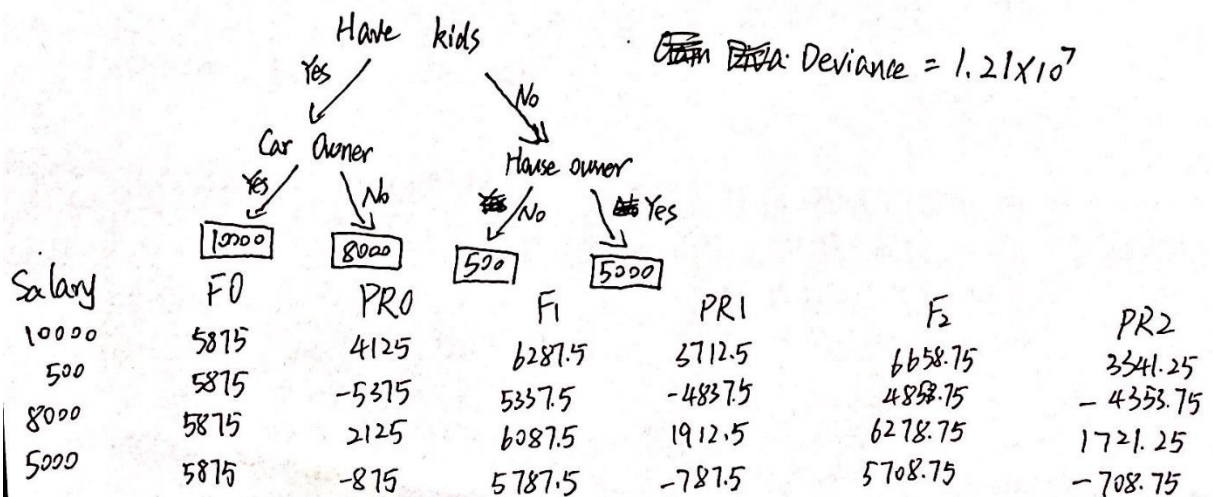
$$= \begin{bmatrix} (\theta_0 + \theta_1 x_1^{(1)} + \dots + \theta_n x_n^{(1)}) + \dots + (\theta_0 + \theta_1 x_1^{(m)} + \dots + \theta_n x_n^{(m)}) \\ x_1^{(1)} \cdot (\theta_0 + \theta_1 x_1^{(1)} + \dots + \theta_n x_n^{(1)}) + \dots + x_1^{(m)} \cdot (\theta_0 + \theta_1 x_1^{(m)} + \dots + \theta_n x_n^{(m)}) \\ \vdots \\ x_n^{(1)} \cdot (\theta_0 + \theta_1 x_1^{(1)} + \dots + \theta_n x_n^{(1)}) + \dots + x_n^{(m)} \cdot (\theta_0 + \theta_1 x_1^{(m)} + \dots + \theta_n x_n^{(m)}) \end{bmatrix}$$

Therefore, $\frac{\partial \theta^T X^T X \theta}{\partial \theta} = 2 X^T X \theta$

2.

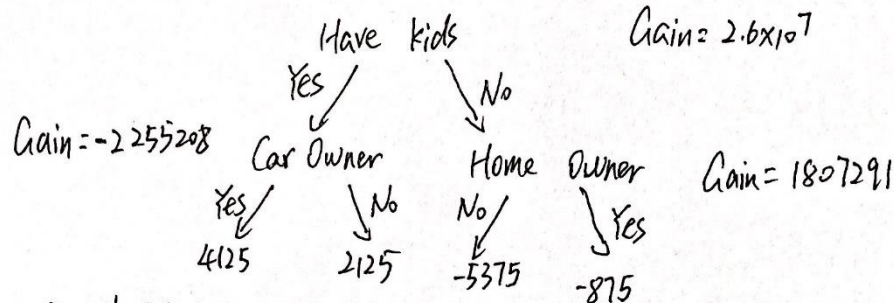
(1)

The tree should be

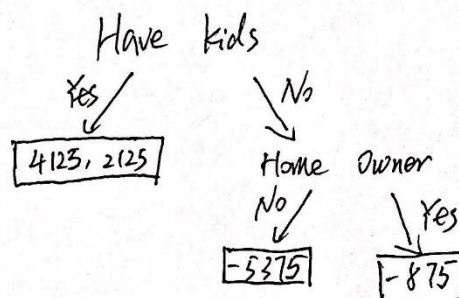


(2)

When I implement the tree I see that



Therefore, it should be pruned.



Salary	F1	PR1	F2	PR2
10000	6083.33	3916.67	6277.78	3722.2
500	5606.25	-5106.25	5350.9	-4852.9
8000	6083.33	1916.67	6277.78	1722.2
5000	5831.25	-831.25	5789.7	-789.7

3.

(1)

For Linear Regression, the result is:

R2 sq: 0.6083953032289722

Mean squared error: 0.53

Test Variance score: 0.60

And the coefficients of features are (the input has been normalized):

HouseAge Latitude AveOccup Longitude Population MedInc AveRooms AveBedrms

model.coef_

```
array([[ 0.47803918, -3.99331079, -4.33923437, -4.40246164,
        -0.19455488,  6.33030675, -15.47310652,  21.11153934]])
```

We can see that House age and Population weight least and AveBedrms weights most. This is probably because that the houses in CA are built during the same time period, and the

public traffic there is very convenient.

(2)

For Gradient Boosting Machine, the result is:

```
R2 sq: 0.8391271074833364
Mean squared error: 0.25
Test Variance score: 0.81
```

And the importance of features is:

```
AveOccup Latitude Longitude AveBedrms Population HouseAge MedInc AveRooms
model.feature_importances_
array([0.13055554, 0.09604826, 0.10855568, 0.0074221 , 0.00636641,
       0.04285832, 0.58718297, 0.02101071])
```

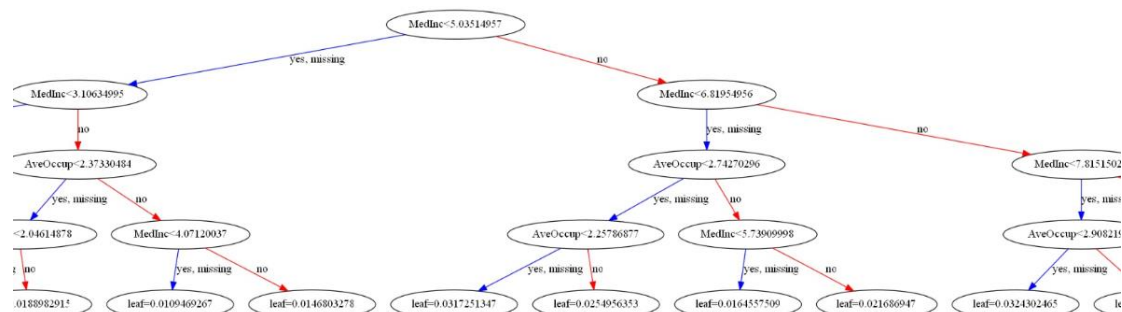
It shows that Median income weights most, which is much more than others. This is reasonable because the median income decides the consumption level, which decides the house price.

(3)

For XGBoosting, the result is:

```
Mean squared error: 0.30
Test Variance score: 0.77
```

The decision tree is:



The importance is:

```
HouseAge AveBedrms MedInc Latitude AveOccup AveRooms Longitude Population
xg_reg.feature_importances_
array([0.06191953, 0.04073134, 0.40880224, 0.09781619, 0.15168677,
       0.13000618, 0.09757712, 0.01146065], dtype=float32)
```

From the importance, we see that it mainly depends on MedInc, AveOccup and AveRooms. They are the only three factors which appear in the decision tree. This model is a little much simple than the GBM, so that it is more stable (lower in variance) but also a little inaccurate (higher in MSE).

APPENDIX

Python Cod:

Linear Regression:

```
#!/usr/bin/env python
```

```
# coding: utf-8
```

```
# In[1]:
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import pandas as pd
```

```
from sklearn import linear_model
```

```
from sklearn.model_selection import train_test_split
```

```
# In[2]:
```

```
dataset = pd.read_csv('C:/Users/69009/Desktop/last/441/problem/CALI.csv')
print(dataset)
```

```
# In[3]:
```

```
feature_names={"MedInc", "HouseAge", "AveRooms", "AveBedrms", "Population",  
               "AveOccup", "Latitude", "Longitude"}
```

```
df_x = pd.DataFrame(dataset, columns = feature_names)
```

```
df_y = pd.DataFrame(dataset, columns = {"HOUSING PRICE"})
```

```
df_x
```

```
# In[4]:
```

```
for column in df_x.columns:
```

```
    df_x[column] = (df_x[column] - df_x[column].min()) / (df_x[column].  
max() - df_x[column].min())
```

```
df_x
```

```
# In[5]:
```

```
x_train, x_test, y_train, y_test = train_test_split(df_x, df_y, test_size = 0.2, random_state = 4)
```

```
# In[6]:
```

```
model = linear_model.LinearRegression()  
model.fit(x_train,y_train)  
results = model.predict(x_test)  
print(results)
```

```
# In[7]:
```

```
model.coef_
```

```
# In[8]:
```

```
from sklearn.metrics import mean_squared_error, r2_score  
model_score = model.score(x_train,y_train)  
print('R2 sq: ', model_score)  
# The mean squared error  
print("Mean squared error: %.2f"% mean_squared_error(y_test, results))  
# Explained variance score: 1 is perfect prediction  
print('Test Variance score: %.2f' % r2_score(y_test, results))
```

```
# In[9]:
```

```
print(np.c_[y_test.values, results][0:5,:])
```

```
# In[ ]:
```

GBM:

```
#!/usr/bin/env python  
# coding: utf-8
```

```
# In[1]:
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn import ensemble
from sklearn import tree
```

```
# In[2]:
```

```
dataset = pd.read_csv('C:/Users/69009/Desktop/last/441/problem/CALI.csv')
feature_names={"MedInc","HouseAge","AveRooms","AveBedrms","Population",
"AveOccup","Latitude","Longitude"}
df_x = pd.DataFrame(dataset, columns = feature_names)
df_y = pd.DataFrame(dataset, columns = {"HOUSING PRICE"})
for column in df_x.columns:
    df_x[column] = (df_x[column] - df_x[column].min()) / (df_x[column].
max() - df_x[column].min())
x_train, x_test, y_train, y_test = train_test_split(df_x, df_y, test_si
ze = 0.2, random_state = 4)
df_x
```

```
# In[3]:
```

```
params = {'n_estimators': 1000, 'max_depth': 4, 'min_samples_split': 2,
'learning_rate': 0.01, 'loss': 'ls'}
model = ensemble.GradientBoostingRegressor(**params)

model.fit(x_train, y_train)
```

```
# In[4]:
```

```
model.feature_importances_
```

```
# In[5]:
```

```

from sklearn.metrics import mean_squared_error, r2_score
model_score = model.score(x_train,y_train)
print('R2 sq: ',model_score)

y_predicted = model.predict(x_test)
# The mean squared error
print("Mean squared error: %.2f"% mean_squared_error(y_test, y_predicte
d))
# Explained variance score: 1 is perfect prediction
print('Test Variance score: %.2f' % r2_score(y_test, y_predicted))

```

In[6]:

```

from sklearn.model_selection import cross_val_predict

fig, ax = plt.subplots()
ax.scatter(y_test, y_predicted, edgecolors=(0, 0, 0))
ax.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k-
-', lw=4)
ax.set_xlabel('Actual')
ax.set_ylabel('Predicted')
ax.set_title("Ground Truth vs Predicted")
plt.show()

```

In[]:

XBG:

```

#!/usr/bin/env python
# coding: utf-8

```

In[1]:

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import mean_squared_error, r2_score
```

```
from sklearn import ensemble
```

```
import xgboost as xgb
```

```
# In[2]:
```

```
dataset = pd.read_csv('C:/Users/69009/Desktop/last/441/problem/CALI.csv')
feature_names={"MedInc","HouseAge","AveRooms","AveBedrms","Population",
"AveOccup","Latitude","Longitude"}
df_x = pd.DataFrame(dataset, columns = feature_names)
df_y = pd.DataFrame(dataset, columns = {"HOUSING PRICE"})
for column in df_x.columns:
    df_x[column] = (df_x[column] - df_x[column].min()) / (df_x[column].
max() - df_x[column].min())
df_x
```

```
# In[3]:
```

```
x_train, x_test, y_train, y_test = train_test_split(df_x, df_y, test_si
ze = 0.2, random_state = 4)
```

```
# In[4]:
```

```
params = {'n_estimators': 500, "objective":"reg:linear",'colsample_bytr
ee': 0.5,'learning_rate': 0.01,
          'max_depth': 4, 'alpha': 1}
xg_reg = xgb.XGBRegressor(**params)

xg_reg.fit(x_train,y_train)
y_predicted = xg_reg.predict(x_test)
```

```
# In[5]:
```

```
xg_reg.feature_importances_
```



```
# In[6]:
```

```
print("Mean squared error: %.2f"% mean_squared_error(y_test, y_predicte  
d))  
# Explained variance score: 1 is perfect prediction  
print('Test Variance score: %.2f' % r2_score(y_test, y_predicted))
```

```
# In[8]:
```

```
params = {"objective": "reg:linear", 'colsample_bytree': 0.5, 'learning_ra  
te': 0.01, 'max_depth': 4, 'alpha': 1}  
data_dmatrix = xgb.DMatrix(data=df_x, label=df_y)  
xg_reg = xgb.train(params=params, dtrain=data_dmatrix, num_boost_round=  
10)  
xgb.plot_tree(xg_reg, num_trees=9)  
plt.rcParams['figure.figsize'] = (70.0, 70.0)  
plt.savefig('C:/Users/69009/Desktop/last/441/problem/XGB/fig1.jpg')  
plt.show()
```

```
# In[ ]:
```